

Network Situational Awareness with tcpdump

Thomas Phillips

Why tcpdump?

Because it's everywhere!
/usr/sbin/tcpdump
(and WinDump, too)




<http://cdn.morguefile.com/imageData/public/files/e/Elfaramawy/10/I/1414048297276uo.jpg>

Topics

The Network Stack

tcpdump Fundamentals

Interesting Recipes



The Network Stack

The Essential Protocols

Ethernet

ARP

Internet Protocol


ICMP

TCP

UDP


Ethernet

https://en.wikipedia.org/wiki/IEEE_802.1Q



Destination and Source MAC Addresses

802.1Q Header (VLAN)



EtherType (<https://www.iana.org/assignments/ieee-802-numbers/ieee-802-numbers.xhtml>)

Payload

Address Resolution Protocol (ARP)

Hardware and Protocol Types

Sender and Target Hardware Addresses

Sender and Target Protocol Addresses

| Internet Protocol (IPv4) over Ethernet ARP packet | | |
|---|---|--------------------------------|
| octet offset | 0 | 1 |
| 0 | Hardware type (HTYPE) | |
| 2 | Protocol type (PTYPE) | |
| 4 | Hardware address length (HLEN) | Protocol address length (PLEN) |
| 6 | Operation (OPER) | |
| 8 | Sender hardware address (SHA) (first 2 bytes) | |
| 10 | (next 2 bytes) | |
| 12 | (last 2 bytes) | |
| 14 | Sender protocol address (SPA) (first 2 bytes) | |
| 16 | (last 2 bytes) | |
| 18 | Target hardware address (THA) (first 2 bytes) | |
| 20 | (next 2 bytes) | |
| 22 | (last 2 bytes) | |
| 24 | Target protocol address (TPA) (first 2 bytes) | |
| 26 | (last 2 bytes) | |

https://en.wikipedia.org/wiki/Address_Resolution_Protocol

Internet Protocol (IPv4)

| IPv4 Header Format | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|--------------------|-------|------------------------|---|---|---|-----|---|---|---|----------|---|----|----|----|----|-----|----|-----------------|----|-----------------|----|----|----|----|----|----|----|----|----|----|----|----|----|
| Offsets | Octet | 0 | | | | | | | | 1 | | | | | | | | 2 | | | | | | | | 3 | | | | | | | |
| Octet | Bit | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| 0 | 0 | Version | | | | IHL | | | | DSCP | | | | | | ECN | | Total Length | | | | | | | | | | | | | | | |
| 4 | 32 | Identification | | | | | | | | | | | | | | | | Flags | | Fragment Offset | | | | | | | | | | | | | |
| 8 | 64 | Time To Live | | | | | | | | Protocol | | | | | | | | Header Checksum | | | | | | | | | | | | | | | |
| 12 | 96 | Source IP Address | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 16 | 128 | Destination IP Address | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 20 | 160 | Options (if IHL > 5) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

IHL (x 4)

Protocol

Total Length (including header)

Source and Destination Addresses

Time To Live (TTL)

Internet Control Message Protocol (ICMP)

| ICMP Header Format | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|--------------------|-------|----------------|---|---|---|---|---|---|---|------|---|----|----|----|----|----|----|----------|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| Offsets | Octet | 0 | | | | | | | | 1 | | | | | | | | 2 | | | | | | | | 3 | | | | | | | |
| Octet | Bit | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| 0 | 0 | Type | | | | | | | | Code | | | | | | | | Checksum | | | | | | | | | | | | | | | |
| 4 | 32 | Rest of Header | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

https://en.wikipedia.org/wiki/Internet_Control_Message_Protocol

Type 8, Code 0 = Echo Request

Type 0, Code 0 = Echo Reply

Type 3 = Destination Unreachable

Type 11 = Time Exceeded

Transmission Control Protocol (TCP)

| TCP Header | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------------|-------|--|---|---|---|-------------------|---|---|--------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-----------------------------|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| Offsets | Octet | 0 | | | | | | | | 1 | | | | | | | | 2 | | | | | | | | 3 | | | | | | | |
| Octet | Bit | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| 0 | 0 | Source port | | | | | | | | | | | | | | | | Destination port | | | | | | | | | | | | | | | |
| 4 | 32 | Sequence number | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 8 | 64 | Acknowledgment number (if ACK set) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 12 | 96 | Data offset | | | | Reserved 0 0 0 | | | N S | C W R | E C E | U R G | A C K | P S H | R S T | S S Y | F I N | Window Size | | | | | | | | | | | | | | | |
| 16 | 128 | Checksum | | | | | | | | | | | | | | | | Urgent pointer (if URG set) | | | | | | | | | | | | | | | |
| 20 | 160 | Options (if data offset > 5. Padded at the end with "0" bytes if necessary.) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| ... | ... | ... | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

https://en.wikipedia.org/wiki/Transmission_Control_Protocol

Source and Destination Ports

Data Offset (x 4)

Flags (SYN, FIN, RST, PSH)

User Datagram Protocol (UDP)

| UDP Header | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|------------|-------|-------------|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|------------------|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| Offsets | Octet | 0 | | | | | | | | 1 | | | | | | | | 2 | | | | | | | | 3 | | | | | | | |
| Octet | Bit | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 |
| 0 | 0 | Source port | | | | | | | | | | | | | | | | Destination port | | | | | | | | | | | | | | | |
| 4 | 32 | Length | | | | | | | | | | | | | | | | Checksum | | | | | | | | | | | | | | | |

https://en.wikipedia.org/wiki/User_Datagram_Protocol

Source and Destination Ports

Length (including header)

tcpdump Fundamentals

The Essential Options



http://cdn.morguefile.com/imageData/public/files/r/ronnieb/preview/fldr_2008_03_13/file0002113164146.jpg

The Basics

Be promiscuous: `ifconfig eth0 promisc`

Filter for addresses, if applicable.

Run as root: `sudo tcpdump`

Filter for ports, if applicable.

Useful options: `-F, -i, -n, -nn, -q, -r, -s, -v, -vv, -vvv, -w`

Filter for specific fields.

Pick a protocol: `ether, arp, ip, icmp, tcp, udp` (there are others, too)

Operations:

`(,),`
`and, or, not,`
`=, !=, <, <=, >, >=,`
`&, |, <<, >>, +, -`

Useful options

-F *file* Read the filter expression from the file.

-i *iface* Listen on a specific interface.

-n Do not do DNS lookups.

-nn Do not use port labels.

-q Quiet mode.

-r *file* Read from a pcap file.

-s *snaplen* Snarf snaplen bytes per packet.

-v, -vv, -vvv Levels of verbosity.

-w *file* Write to a pcap file.

Core Filters

Example: `sudo tcpdump arp src host 192.168.1.1`

Protocol:

ether

arp

ip

icmp

tcp

udp

Address:

host *address*

src host *address*

dst host *address*

Port:

port *num*

src port *num*

dst port *num*

Filter for Specific Fields

protocol[start:length] op value

IPv4 Time To Live

`ip[8]`

TTL is less than 3

TTL `ip[8] < 3`

IPv4 Total Length

`ip[2:2]`

IPv4 Header Length

`(ip[0] & 0xf) << 2`

TCP Data Offset

`(tcp[12] & 0xf0) >> 2`

TCP segments that contain data

(
 (
 Total Length **Header Length**
 (`ip[2:2] - (ip[0] & 0xf) << 2`) -
 (`(tcp[12] & 0xf0) >> 2`)
)
 Data Offset
 != 0
)

Interesting Recipes

Good Stuff



<http://cdn.morguefile.com/imageData/public/files/m/Moonlightway/10//1414561609wqz2o.jpg>

Interesting Recipes

Traffic not on my VLAN (data leakage)

Traceroute activity

Fragmented IP traffic

TCP traffic leaving a subnet

Unauthorized TCP ports

List of endpoints calling this endpoint (TCP, UDP)

List of endpoints this endpoint is calling out to (TCP, UDP)

HTTP POST requests

Traffic not on my VLAN (data leakage)

```
tcpdump \  
-vv \  
-i eth0 \  
' ( vlan and ( ether[14:2] & 0xffff != 1000 ) ) '
```

Look for 802.1Q frames tagged with VLAN

Adapted from:

<http://serverfault.com/questions/196250/tcpdump-capture-one-of-several-vlans>

Traceroute activity

```
tcpdump \  
-vv \  
-i eth0 \  
' ( ip[8] < 3 ) '
```

Low TTL

Protocol other than IPv4 is not important.

Fragmented IP traffic

```
tcpdump \  
-vv \  
-i eth0 \  
' ( (ip[6] & 64 == 0) and (ip[6:2] > 0) ) '
```

Don't fragment bit is not set.

Fragment Offset

Adapted from:

<https://blog.wains.be/2007/2007-10-01-tcpdump-advanced-filters.md>

TCP traffic leaving a subnet

```
tcpdump \
```

```
-vv \
```

```
-i eth0 \
```

```
'(tcp[13] & 18 == 2) and not dst net 192.168.1.0/24'
```

SYN flag set, no ACK

Our subnet

Unauthorized TCP ports

```
tcpdump \
```

```
-vv \
```

```
-i eth0 \
```

Our IP address

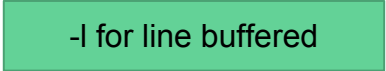
```
`dst host 192.168.1.100 and \
```

```
(tcp[13] & 18 == 2) and (not dst port 22)'
```

SYN flag set, no ACK

Display anything not SSH.

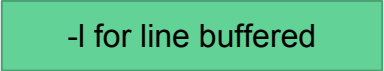
List of endpoints calling this endpoint (TCP, UDP)

```
tcpdump \    
-i eth0 -l \   
'(udp or (tcp[13] & 18 == 2)) and \   
dst host 192.168.1.100 | cut -d' ' -f3,6-
```

Example output:

```
zambonia.domain 18870* 1/0/0 PTR Zambonia. (64)   
qb-in-f189.1e100.net.https UDP, length 41   
qb-in-f189.1e100.net.https UDP, length 20   
qh-in-f189.1e100.net.https UDP, length 36   
qb-in-f189.1e100.net.https UDP, length 33   
qb-in-f189.1e100.net.https UDP, length 63   
qb-in-f189.1e100.net.https UDP, length 39   
zambonia.domain 65529 2/0/0 CNAME fd-fp3.wg1.b.yahoo.com., AAAA 2001:4998:58:c02::a9 (86)
```

List of endpoints this endpoint is calling out to (TCP, UDP)

```
tcpdump \  -l for line buffered  
-i eth0 -1 -nn \  
'(udp or (tcp[13] & 18 == 2)) and \  
not dst net 192.168.1.0/24' | cut -d' ' -f5-
```

Example output:

```
74.125.22.189.443: UDP, length 38  
98.139.180.149.80: Flags [S], seq 1112211606, win 65535, options [mss 1460,nop,wscale 5,  
nop,nop,TS val 440461606 ecr 0,sackOK,eol], length 0  
172.217.1.195.443: UDP, length 1350  
172.217.1.195.443: UDP, length 337  
172.217.1.195.443: Flags [S], seq 640180255, win 65535, options [mss 1460,nop,wscale 5,nop,  
nop,TS val 440462066 ecr 0,sackOK,eol], length 0  
172.217.1.195.443: UDP, length 37  
172.217.1.195.443: UDP, length 40
```


HTTP POST requests

http://cdn.morguefile.com/imageData/public/files/s/seeka/preview/fldr_2004_04_14/file0001535568481.jpg

```
tcpdump \
  -i eth0 -l \
  'tcp port 80 and \
    (((ip[2:2] - ((ip[0]&0xf)<<2)) - ((tcp[12]&0xf0)>>2)) !=
0)' | egrep --line-buffered 'POST'
```

Adapted from:

<https://blog.wains.be/2007/2007-10-01-tcpdump-advanced-filters.md>

Example output:

```
23:51:13.012659 IP dunwich.hsd1.md.comcast.net.52660 > us-3-zone-1.syronex.com.http: Flags
[P.], seq 1451860020:1451860913, ack 4079266449, win 4117, options [nop,nop,TS val
442006656 ecr 356004271], length 893: HTTP: POST /form/yx/submit HTTP/1.1
```

Searching for the POST
string in the filter is left as an
exercise for the reader.



```
bash-3.2# echo -n "POST" | od -t x1c
0000000  50  4f  53  54
          P  O  S  T
```

El fin.